

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 import numpy as np
5 sns.set_theme(color_codes = True)
```

Dataset :

<https://www.kaggle.com/datasets/abhia1999/chronic-kidney-disease>


```
1 df = pd.read_csv('new_model.csv')
2 df
```

	Bp	Sg	Al	Su	Rbc	Bu	Sc	Sod	Pot	Hemo	Wbcc	Rbcc	Htn	Class
0	80.0	1.020	1.0	0.0	1.0	36.0	1.2	137.53	4.63	15.4	7800.0	5.20	1.0	1
1	50.0	1.020	4.0	0.0	1.0	18.0	0.8	137.53	4.63	11.3	6000.0	4.71	0.0	1
2	80.0	1.010	2.0	3.0	1.0	53.0	1.8	137.53	4.63	9.6	7500.0	4.71	0.0	1
3	70.0	1.005	4.0	0.0	1.0	56.0	3.8	111.00	2.50	11.2	6700.0	3.90	1.0	1
4	80.0	1.010	2.0	0.0	1.0	26.0	1.4	137.53	4.63	11.6	7300.0	4.60	0.0	1
...
395	80.0	1.020	0.0	0.0	1.0	49.0	0.5	150.00	4.90	15.7	6700.0	4.90	0.0	0
396	70.0	1.025	0.0	0.0	1.0	31.0	1.2	141.00	3.50	16.5	7800.0	6.20	0.0	0
397	80.0	1.020	0.0	0.0	1.0	26.0	0.6	137.00	4.40	15.8	6600.0	5.40	0.0	0
398	60.0	1.025	0.0	0.0	1.0	50.0	1.0	135.00	4.90	14.2	7200.0	5.90	0.0	0
399	80.0	1.025	0.0	0.0	1.0	18.0	1.1	141.00	3.50	15.8	6800.0	6.10	0.0	0

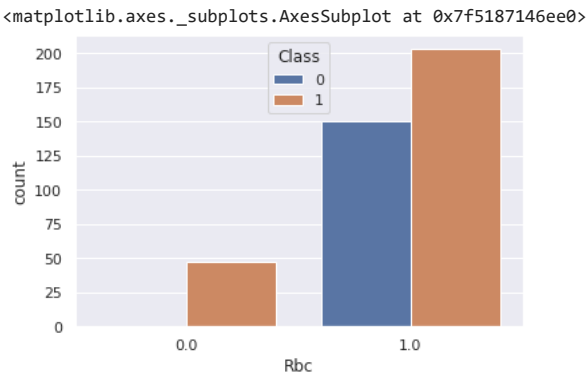
400 rows × 14 columns

Exploratory Data Analysis


```
1 sns.countplot(data=df, x="Htn", hue="Class")
```



```
1 sns.countplot(data=df, x="Rbc", hue="Class")
```



```
1 sns.histplot(data=df, x="Bp", hue="Class", multiple="stack")
```



▼ Data Preprocessing

```
1 df.isnull().sum()

Bp      0
Sg      0
Al      0
Su      0
Rbc     0
Bu      0
Sc      0
Sod     0
Pot     0
Hemo    0
Wbcc    0
Rbcc    0
Htn     0
Class   0
dtype: int64

1 #replace 0 value with NaN
2 df_copy = df.copy(deep = True) #deep = True -> Buat salinan indeks dan data dalam dataframe
3 df_copy[['Bp', 'Sg', 'Bu', 'Sc', 'Sod', 'Pot', 'Hemo', 'Wbcc', 'Rbcc']] = df_copy[['Bp', 'Sg', 'Bu', 'Sc', 'Sod', 'Pot', 'Hemo', 'Wbcc', 'Rbcc']].replace(
4
5 # Showing the Count of NaNs
6 print(df_copy.isnull().sum())

Bp      0
Sg      0
Al      0
Su      0
Rbc     0
Bu      0
Sc      0
Sod     0
Pot     0
Hemo    0
Wbcc    0
Rbcc    0
Htn     0
Class   0
dtype: int64
```

▼ Check if the class label is balanced or not




▼ Do Oversampling Minority Class to Balance the class label

```
1 from sklearn.utils import resample
2 #create two different dataframe of majority and minority class
3 df_majority = df[(df['Class']==1)]
4 df_minority = df[(df['Class']==0)]
5 # upsample minority class
6 df_minority_upsampled = resample(df_minority,
7                                 n_samples= 250,
8                                 random_state=0)
9
10 # Combine majority class with upsampled minority class
11 df2 = pd.concat([df_minority_upsampled, df_majority])


1 sns.countplot(df2['Class'])
2 print(df2.Class.value_counts())
```

```
/usr/local/lib/python3.8/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From
warnings.warn(
0    250
1    250
Name: Class, dtype: int64
```




▼ Check the Outlier using Boxplot


```
1 sns.boxplot(x=df2["Bp"])
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186ed9970>
```




```
1 sns.boxplot(x=df2["Sg"])
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186e83d90>
```



```
1 sns.boxplot(x=df2["Bu"])
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186de2be0>
```




```
1 sns.boxplot(x=df2["Sc"])
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186dc1dc0>
```




```
1 sns.boxplot(x=df2["Sod"])
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186d99970>
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186d0c790>
```



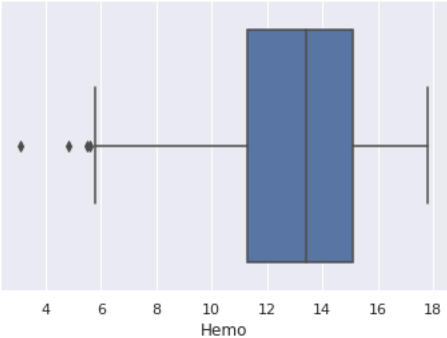
```
1 sns.boxplot(x=df2["Pot"])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186e8aa60>
```




```
1 sns.boxplot(x=df2["Hemo"])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186cc4220>
```




```
1 sns.boxplot(x=df2["Wbcc"])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186c775b0>
```



```
1 sns.boxplot(x=df2["Rbcc"])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5186be32b0>
```



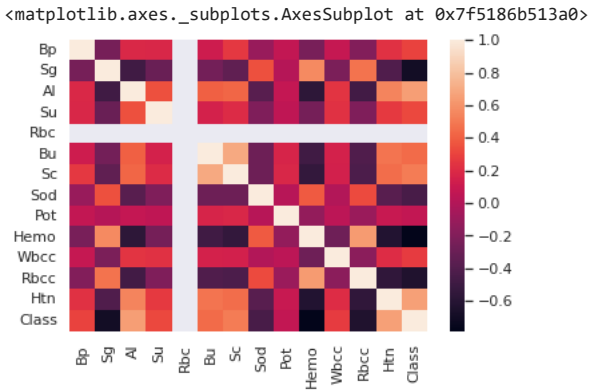
▼ Remove Outlier using Z-Score

```
1 import scipy.stats as stats
2 z = np.abs(stats.zscore(df2))
3 data_clean = df2[(z<3).all(axis = 1)]
4 data_clean.shape
```

```
(420, 14)
```


Heatmap Data Correlation

```
1 sns.heatmap(data_clean.corr(), fmt='.2g')
```



```
1 #Rbc attribute is irrlevant, so we have to remove it
2 data_clean2 = df.drop(columns=['Rbc'])

1 corr = data_clean2[data_clean2.columns[1:]].corr()['Class'][:-1]
2 plt.plot(corr)
3 plt.xticks(rotation=90)
4 plt.show()
```



Machine Learning Model Building

```
1 X = data_clean2.drop('Class', axis=1)
2 y = data_clean2['Class']
```

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.metrics import accuracy_score
3 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,random_state=0)
```

Random Forest

```
1 from sklearn.ensemble import RandomForestClassifier
2 rfc = RandomForestClassifier(random_state=0)
3 rfc.fit(X_train, y_train)
```

RandomForestClassifier(random_state=0)


```
1 y_pred = rfc.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 100.0 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 1.0
Precision Score : 1.0
Recall Score : 1.0

```
1 from sklearn.metrics import classification_report, confusion_matrix
2 cm = confusion_matrix(y_test, y_pred)
3 plt.figure(figsize=(5,5))
4 sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')
5 plt.ylabel('Actual label')
6 plt.xlabel('Predicted label')
7 all_sample_title = 'Accuracy Score: {0}'.format(rfc.score(X_test, y_test)*100)
8 plt.title(all_sample_title, size = 15)
```



▼ KNearest Neighbor

```
1 from sklearn.neighbors import KNeighborsClassifier
2 knn = KNeighborsClassifier()
3 knn.fit(X_train, y_train)
```

KNeighborsClassifier()


```
1 y_pred = knn.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 62.5 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.6875
Precision Score : 0.75
Recall Score : 0.6346153846153846

```
1 from sklearn.metrics import classification_report, confusion_matrix
2 cm = confusion_matrix(y_test, y_pred)
3 plt.figure(figsize=(5,5))
4 sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')
5 plt.ylabel('Actual label')
6 plt.xlabel('Predicted label')
7 all_sample_title = 'Accuracy Score: {0}'.format(knn.score(X_test, y_test)*100)
8 plt.title(all_sample_title, size = 15)
```



▼ AdaBoost

```
1 from sklearn.ensemble import AdaBoostClassifier
2 ada = AdaBoostClassifier(random_state=0)
3 ada.fit(X_train, y_train)
```

AdaBoostClassifier(random_state=0)

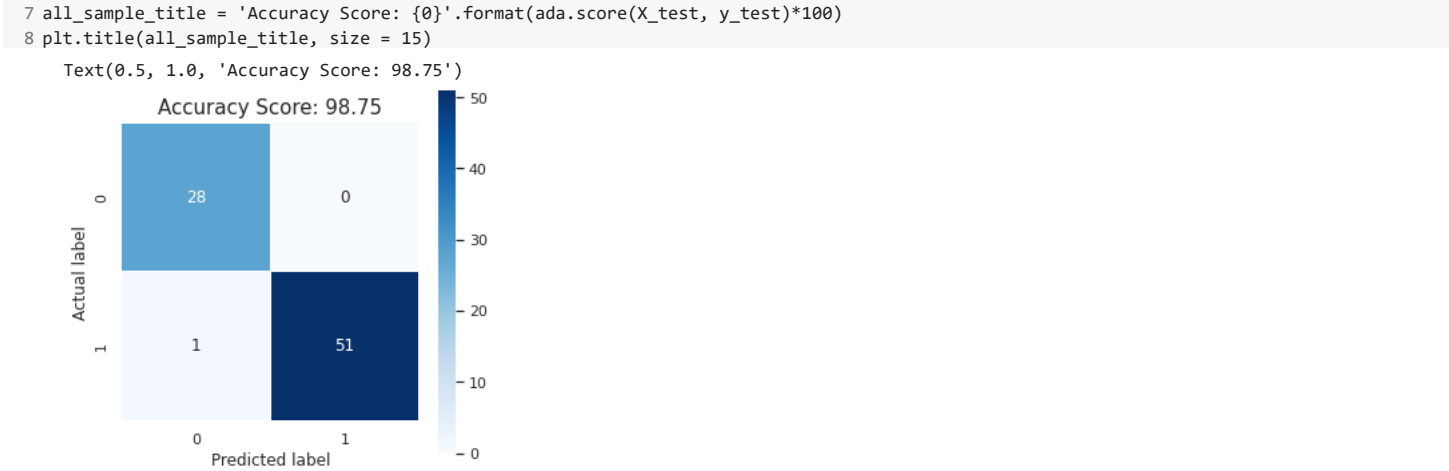
```
1 y_pred = ada.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 98.75 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.9902912621359222
Precision Score : 1.0
Recall Score : 0.9807692307692307

```
1 from sklearn.metrics import classification_report, confusion_matrix
2 cm = confusion_matrix(y_test, y_pred)
3 plt.figure(figsize=(5,5))
4 sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap = 'Blues')
5 plt.ylabel('Actual label')
6 plt.xlabel('Predicted label')
```



▼ Logistic Regression

```
1 from sklearn.linear_model import LogisticRegression
2 lr = LogisticRegression(random_state = 0)
3 lr.fit(X_train, y_train)

/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

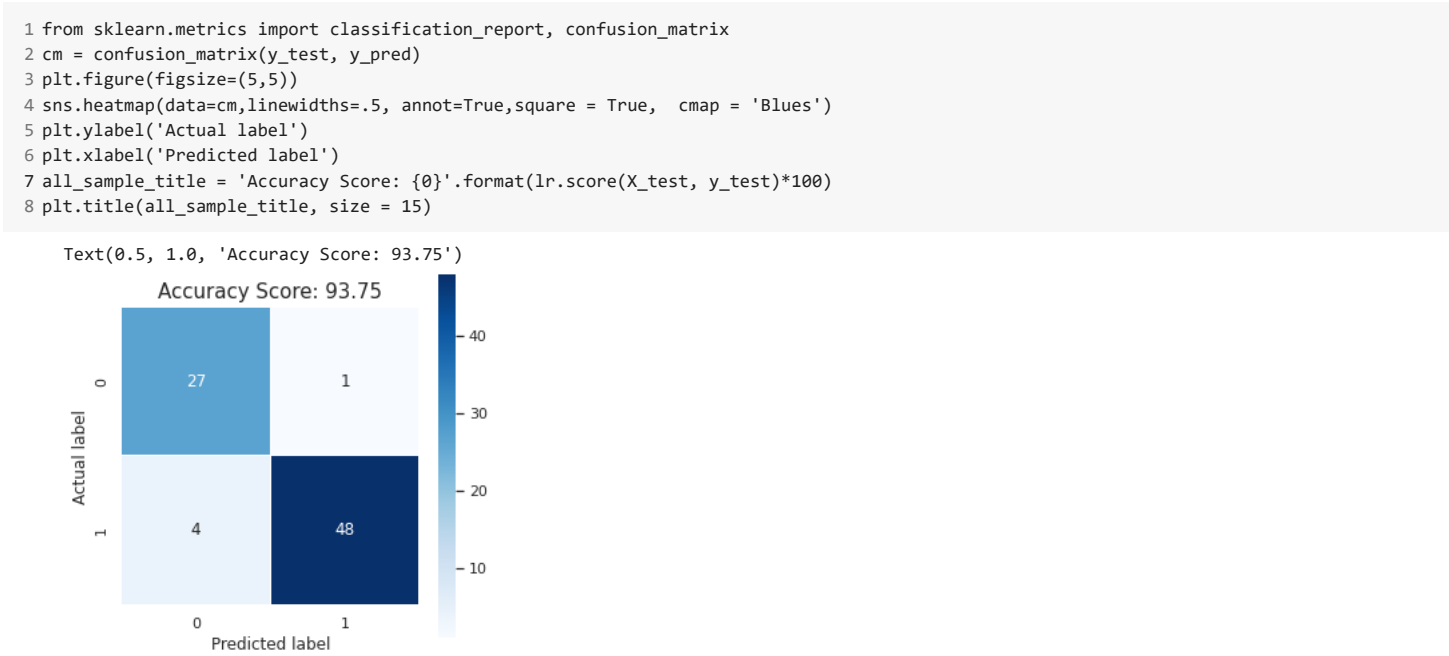
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(
LogisticRegression(random_state=0)

1 y_pred = lr.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")

Accuracy Score : 93.75 %
```

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))

F-1 Score : 0.9504950495049506
Precision Score : 0.9795918367346939
Recall Score : 0.9230769230769231
```



✓ 0s completed at 1:05 PM

● ×